

Supplementary Material: Bayesian estimation in scale mixture of skew-normal linear mixed models using Hamiltonian Monte Carlo

Fernanda L Schumacher, Larissa A Matos, Pedro L Ramos, and Francisco Louzada

Contents

Introduction	1
Loading packages	1
Data Generation	2
AR(1)-ST-LMM Fitting using Stan	3
Model comparison with a simpler, Gaussian model	7
Sensitivity analysis: changing prior distribution for shape parameter	9
Frequentist Comparison (EM Algorithm)	10

Introduction

This document provides the full reproducible workflow for the simulation and estimation of the AR(1) Skew-t Linear Mixed Model (AR1-ST-LMM) discussed in the main text, in addition for comparisons with a Gaussian model and a brief sensitivity analysis for the degrees of freedom parameter ν . All materials needed to reproduce this code are available at <https://github.com/fernandalschumacher/HMC-SMSN-LMM>.

Stan is required for running this code. For information on installing Stan, please check <https://github.com/stan-dev/rstan/wiki/RStan-Getting-Started#installation-of-rstan>.

Loading packages

```
library(tictoc)
library(dplyr)
library(tidyr)
library(purrr)
library(ggplot2)
library(loo)
library(rstan)
library(posterior)
library(bayesplot)
library(HDIInterval)
library(skewlmm)
library(bayesplot)
```

Data Generation

We generate a longitudinal dataset with $m = 200$ subjects and $n_j = 5$ observations per subject from an AR(1)-ST-LMM using the `skewlmm` package.


```
set.seed(955)
nj1 <- 5
m <- 200
D1_true <- matrix(c(.5,.1,.1,1), ncol=2)

gendatList <- map(rep(nj1, m), function(nj) {
  skewlmm::rsmn.lmm(1:nj, cbind(1, 1:nj), cbind(1, 1:nj),
    sigma2 = .25, D1 = D1_true, beta = c(1, 2),
    lambda = c(1,3), depStruct = "ARp",
    phi = .5, distr = 'st', nu = 6)
})

gendat <- bind_rows(gendatList, .id = "ind") %>%
  mutate(indnum = as.numeric(ind)) %>%
  arrange(indnum)

# Visualization
ggplot(gendat, aes(x = x, y = y, group = ind)) +
  geom_line(alpha = 0.3) +
  stat_summary(aes(group = 1), geom = "line", fun = mean, col = "blue", size = 1.5) +
  theme_minimal() +
  labs(title = "Simulated Individual Trajectories", x = "Time", y = "Response")
```

Simulated Individual Trajectories



AR(1)-ST-LMM Fitting using Stan

Setup Data for Stan

```
N <- nrow(gendat)
n <- n_distinct(gendat$ind)
x_mat <- model.matrix(y ~ x, data = gendat)
z_mat <- model.matrix(~ x, data = gendat)
l_cols <- ncol(x_mat)
q1_cols <- ncol(z_mat)
njvec <- as.numeric(table(gendat$indnum))

data_list_stan <- list(
  y = gendat$y, x = x_mat, z = z_mat,
  timevar = gendat$time, indnum = gendat$indnum
)
```

Fitting the ST Model

```
tic("AR(1)-ST-LMM Fit")
tbayes <- stan(file = 'lmm-AR1-ST.stan',
  data = list(N=N, n=n, l=l_cols, q1=q1_cols, njvec=njvec,
```

```

        y=gendat$y, x=x_mat, z=z_mat,
        timevar=gendat$time, ind=gendat$indnum, sdLP=2),
        thin = 5, chains = 3, iter = 5000, warmup = 1000, seed = 9955)
toc()


## AR(1)-ST-LMM Fit: 1837.29 sec elapsed

print(tbayes, pars=c("beta","sigmae","phi1","D1","lambda","nu"), probs=c(.025,.975), digits=3)

## Inference for Stan model: anon_model.
## 3 chains, each with iter=5000; warmup=1000; thin=5;
## post-warmup draws per chain=800, total post-warmup draws=2400.
##
##          mean se_mean    sd   2.5% 97.5% n_eff  Rhat
## beta[1]   0.898   0.001 0.068   0.765 1.033  2032 1.001
## beta[2]   2.081   0.001 0.060   1.973 2.207  2200 1.002
## sigmae    0.528   0.001 0.024   0.478 0.575   996 1.002
## phi1      0.581   0.003 0.071   0.415 0.696   583 1.007
## D1[1,1]   0.092   0.005 0.100   0.000 0.352   427 1.010
## D1[1,2]   0.029   0.002 0.078  -0.131 0.189  1791 1.000
## D1[2,1]   0.029   0.002 0.078  -0.131 0.189  1791 1.000
## D1[2,2]   0.826   0.005 0.159   0.544 1.139   926 1.003
## lambda[1] -0.068   0.063 0.963  -2.013 1.841   233 1.016
## lambda[2]  2.938   0.032 0.921   1.453 4.956   835 1.001
## nu        5.435   0.022 0.924   3.899 7.480  1744 1.000
##
## Samples were drawn using NUTS(diag_e) at Fri Dec 26 18:04:37 2025.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

mcmc_dens_chains(tbayes, pars = c("beta[1]","beta[2]"))+
  geom_vline(xintercept = 0, linetype = "dotted") + theme_bw()

```



Checking Acceptance Probability

```
sampler_params <- get_sampler_params(tbayes, inc_warmup = FALSE)
mean(unlist(lapply(sampler_params, function(x) x[, "accept_stat__"])))
```

```
## [1] 0.9179794
```

Computing Highest (Posterior) Density Intervals (HDI)


```
tsample <- rstan::extract(tbayes, pars = c("beta[1]", "beta[2]", "sigmae", "phi1", "D1[1,1]",
                                           "D1[1,2]", "D1[2,2]", "lambda[1]", "lambda[2]", "nu"))
HDI <- HDInterval::hdi(tsample) %>% unlist()
data.frame(lower = HDI[c(TRUE, FALSE)], upper = HDI[c(FALSE, TRUE)]) %>% knitr::kable()
```

	lower	upper
beta[1].lower	0.7558937	1.0221253
beta[2].lower	1.9744829	2.2081761
sigmae.lower	0.4747794	0.5712310
phi1.lower	0.4363855	0.7109411
D1[1,1].lower	0.0000000	0.3038586

	lower	upper
D1[1,2].lower	-0.1325639	0.1868039
D1[2,2].lower	0.5440794	1.1395394
lambda[1].lower	-2.0187785	1.8202168
lambda[2].lower	1.3409993	4.8007222
nu.lower	3.7751090	7.2890852

Exploring the Posterior Distributions

```
bayesplot::color_scheme_set("red")
bayesplot::mcmc_rank_hist(tbayes, pars=c("beta[1]", "beta[2]", "phi1", "lambda[1]", "lambda[2]", "nu"))
```



Additional exploration can be easily done by launching the shiny application using `launch_shinystan(tbayes)`.


Computing Information Criteria

To evaluate the marginal log-likelihood function, an approximation of Leave-One-Out Cross-Validation (LOO), and the Widely applicable information criterion (WAIC) for all MCMC samples, we can use the function `criteriaAR1`:

```
source("auxfunctions-bayes.R")
```

```
# ST Model LOO
st_loo <- criteriaAR1(tbayes, data_list_stan, distr = "st")

# Checking the Pareto smoothed importance sampling (PSIS)
plot(st_loo$loo)
```



Model comparison with a simpler, Gaussian model


We compare the ST model against a standard Normal model.

```
# Normal Model Fitting
tic("AR(1)-N-LMM Fit")
nbayes <- stan(file = 'lmm-AR1-N.stan',
  data = list(N=N, n=n, l=l_cols, q1=q1_cols, njvec=njvec,
    y=gendat$y, x=x_mat, z=z_mat,
    timevar=gendat$time, ind=gendat$indnum),
  thin = 5, chains = 3, iter = 5000, warmup = 1000, seed = 9955)
toc()
```

```
## AR(1)-N-LMM Fit: 1509.05 sec elapsed
```

```
# Computing the information criteria
normal_loo <- criteriaAR1(nbayes, data_list_stan, distr = "sn")
```

```
# Checking the Pareto smoothed importance sampling (PSIS)
plot(normal_loo$loo)
```



Comparing the models based on LOO

```
comp_loo <- loo_compare(list(ST = st_loo$loo, Normal = normal_loo$loo))
knitr::kable(comp_loo, caption = "Model Comparison: LOO Criteria", simplify = FALSE)
```

Table 2: Model Comparison: LOO Criteria

	elpd_diff	se_diff	elpd_loo	se_elpd_loo	p_loo	se_p_loo	looic	se_looic
ST	0.00000	0.00000	-1370.738	30.48947	6.212676	0.7299541	2741.476	60.97893
Normal	-82.49983	30.19118	-1453.238	48.94037	8.824400	5.0013610	2906.476	97.88074

Comparing the models based on WAIC

```
comp_loo <- loo_compare(list(ST = st_loo$waic, Normal = normal_loo$waic))
knitr::kable(comp_loo, caption = "Model Comparison: WAIC Criteria", simplify = FALSE)
```


Table 3: Model Comparison: WAIC Criteria

	elpd_diff	se_diff	elpd_waic	se_elpd_waic	p_waic	se_p_waic	waic	se_waic
ST	0.00000	0.00000	-1370.727	30.48860	6.202024	0.7303406	2741.455	60.97721
Normal	-82.48586	30.17448	-1453.213	48.92392	8.799775	4.9826414	2906.426	97.84784

Sensitivity analysis: changing prior distribution for shape parameter

To evaluate the model's sensitivity to the choice of the prior distribution for ν , we now compare our original results based on a truncated Cauchy distribution with the ones considering a Gamma(2, 0.2) prior distribution for ν .

```
tic("AR(1)-ST-LMM Fit with Gamma prior")
tbayes_nu <- stan(file='lmm-AR1-ST-nuprior.stan',
  data = list(N=N, n=n, l=l_cols, q1=q1_cols, njvec=njvec,
    y=gendat$y, x=x_mat, z=z_mat,
    timevar=gendat$time, ind=gendat$indnum, sdLP=2),
  thin = 5, chains = 3, iter = 5000, warmup = 1000, seed = 9955,
  control = list(adapt_delta=.95))
toc()
```

```
## AR(1)-ST-LMM Fit with Gamma prior: 2583.43 sec elapsed
```

```
print(tbayes_nu,par=c("beta","sigmae","phi1","D1","lambda","nu"),
  probs = c(.025,.975), digits=3)
```

```
## Inference for Stan model: anon_model.
## 3 chains, each with iter=5000; warmup=1000; thin=5;
## post-warmup draws per chain=800, total post-warmup draws=2400.
##
##               mean se_mean    sd   2.5% 97.5% n_eff Rhat
## beta[1]      0.897   0.001 0.068   0.764 1.026  2390 1.001
## beta[2]      2.079   0.001 0.058   1.966 2.196  2318 1.000
## sigmae       0.528   0.001 0.026   0.475 0.575   807 1.002
## phi1         0.570   0.004 0.080   0.372 0.690   482 1.004
## D1[1,1]      0.112   0.005 0.116   0.000 0.413   489 1.001
## D1[1,2]      0.031   0.002 0.082  -0.142 0.194  1542 1.000
## D1[2,1]      0.031   0.002 0.082  -0.142 0.194  1542 1.000
## D1[2,2]      0.851   0.004 0.162   0.569 1.206  1344 1.001
## lambda[1]   -0.044   0.041 0.884  -1.834 1.764   457 1.003
## lambda[2]    2.941   0.030 0.925   1.499 5.057   982 1.006
## nu           5.732   0.023 0.991   4.110 7.993  1842 1.002
##
## Samples were drawn using NUTS(diag_e) at Fri Dec 26 19:21:57 2025.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

```
## Comparing with original model
posterior1 <- posterior::as_draws_matrix(tbayes)
posterior2 <- posterior::as_draws_matrix(tbayes_nu)
```


```

library(posterior)
d1 <- as_draws_df(posterior1) |>
  dplyr::select(nu, starts_with("beta")) |>
  tidyr::pivot_longer(cols = everything(), names_to = "param", values_to = "values") |>
  dplyr::mutate(Model = "Cauchy Prior")
d2 <- as_draws_df(posterior2) |>
  dplyr::select(nu, starts_with("beta")) |>
  pivot_longer(cols = everything(), names_to = "param", values_to = "values") |>
  dplyr::mutate(Model = "Gamma Prior")

d <- dplyr::bind_rows(d1, d2)

ggplot(d, aes(x = values, fill = Model)) +
  geom_density(alpha = 0.4) +
  ggtitle("Sensitivity of nu to Prior Choice") +
  theme_minimal() + facet_wrap(~param, scale = "free")

```



Frequentist Comparison (EM Algorithm)

For validation, we compare our Bayesian estimates with the frequentist EM algorithm results.

```

tic("Frequentist Model")
fit_EM <- smsn.lmm(data = gendat, formFixed = y~x, groupVar = 'ind',
  formRandom = ~x, depStruct = "ARp", timeVar = 'time',

```

```
distr = 'st')
toc()
```

```
## Frequentist Model: 79.33 sec elapsed
```

```
summary(fit_EM)
```

```
## Linear mixed models with distribution st and dependency structure ARp
## Call:
## smsn.lmm(data = gendat, formFixed = y ~ x, groupVar = "ind",
##   formRandom = ~x, depStruct = "ARp", timeVar = "time", distr = "st")
##
## Distribution st with nu = 5.494048
##
## Random effects:
##   Formula: ~x
##   Structure:
##   Estimated variance (D):
##           (Intercept)           x
## (Intercept) 0.13344526 0.04799053
## x           0.04799053 0.85415275
##
## Fixed effects: y ~ x
## with approximate confidence intervals
##           Value Std.error CI 95% lower CI 95% upper
## (Intercept) 0.8998495 0.0712506   0.7602009   1.039498
## x           2.0808533 0.0588821   1.9654465   2.196260
##
## Dependency structure: ARp
##   Estimate(s):
##   sigma2      phi1
## 0.2702294 0.5484190
##
## Skewness parameter estimate: 0.05175053 2.698154
##
## Model selection criteria:
##   logLik      AIC      BIC
## -1324.429 2668.858 2717.936
##
## Number of observations: 1000
## Number of groups: 200
```